Xingjian Diao

Email: xingjian.diao.gr@dartmouth.edu | Website: https://xid32.github.io/

in Linkedin | ○ Github | ⊗ Google Scholar

RESEARCH INTERESTS

My research focuses on **multimodal learning** for video, audio, and language understanding. I have developed methods for **multimodal reasoning**, **efficient multimodal learning**, and **generative multimodal modeling**, aiming to build scalable and generalizable multimodal models that advance multimodal question answering, video understanding, and audio–visual reasoning across complex real-world scenarios and dynamic environments. My GitHub repositories on multimodal large language models (MLLMs) have received 1.5k+ Stars ⭐ .

EDUCATION

- **Dartmouth College** *Sep 2022 - Jun 2026 (Expected)*
  Ph.D. student in Computer Science Hanover, USA
  - **Advisor**: Prof. Soroush Vosoughi and Prof. Jiang Gui

- **Northwestern University** *Sep 2020 - Dec 2021*
  Master of Science, Computer Science Evanston, IL
  - **Advisor**: Prof. Nabil Alshurafa

- **University of Pittsburgh** *Aug 2016 - Apr 2020*
  Bachelor of Science, Computer Science Pittsburgh, PA

INTERNSHIP

- **Amazon** *June 2025 - Sept 2025*
  Applied Scientist Intern Santa Cruz, USA

  High-Frequency Video-to-IMU Synthesis via Physics-Guided Simulation and Hybrid U-Net Refinement

  Pdf | Proposed PrimeIMU, a physics-guided video-to-IMU generation framework that fuses low-frequency kinematic cues from 3D video poses with physics-inspired simulated inertial initialization through a hybrid U-Net refinement module, effectively bridging the anatomical–inertial gap to produce high-fidelity, sensor-faithful IMU signals that generalize across activities, devices, and datasets, enabling synthetic-only training, cross-domain adaptation, and scalable deployment of wearable sensing models.

  🏷 Computer Vision 🏷 Ubiquitous Computing

  IMU2Reason: A Multimodal LLM for Safety-Aware Activity Understanding

  Pdf | Proposed MHIA-LM, a multimodal LLM trained on fused video and IMU representations, where spatiotemporal video embeddings from InternVideo2 and IMU motion features encoded by a PatchTST-based encoder are integrated through a multi-layer Gated Hierarchical Interaction module and mapped into the LLM token embedding space for instruction tuning, enabling fine-grained activity recognition and safety reasoning.

  🏷 MLLM 🏷 Multimodal Reasoning

SELECTED 1ˢᵀ-AUTHOR PUBLICATIONS [FULL LIST]

[EMNLP 2025] SoundMind: RL-Incentivized Logic Reasoning for Audio-Language Models

Oral Presentation Award , 30ᵗʰ EMNLP

Xingjian Diao , Chunhui Zhang, Keyi Kong, Weiyi Wu, Chiyu Ma, Zhongyu Ouyang, Peijun Qing, Soroush Vosoughi, Jiang Gui

Pdf | Github Code; ⭐ **Starred 1k+** | Dataset | Introduced the Audio Logical Reasoning (ALR) task containing 6,446 text–audio CoT-annotated samples to enable complex reasoning over spoken content, and proposed SoundMind, a rule-based reinforcement learning algorithm that enhances deep cross-modal reasoning in audio-language models.

🏷 Large Audio-Language Model 🏷 Reinforcement Learning 🏷 Multimodal Reasoning

[NAACL 2025] Temporal Working Memory: Query-Guided Temporal Segment Refinement for Enhanced Multimodal Understanding

Guarini Graduate Student Travel Award , Dartmouth College

Xingjian Diao , Chunhui Zhang, Weiyi Wu, Zhongyu Ouyang, Peijun Qing, Ming Cheng, Soroush Vosoughi, Jiang Gui

Pdf | Github Code; ⭐ **Starred 300+** | Proposed Temporal Working Memory, a plug-and-play query-guided segment refinement module that maintains dynamic temporal memory to effectively preserve task-relevant video–audio segment and enhance long-range temporal reasoning, achieving consistent performance gains when integrated into nine recent state-of-the-art multimodal large language models (MLLMs) across AVQA, video captioning, and retrieval tasks.

🏷 MLLM 🏷 Video Understanding

[EMNLP 2024] Learning Musical Representations for Music Performance Question Answering

BMDS Travel Award , Dartmouth College

Xingjian Diao , Chunhui Zhang, Tingxuan Wu, Ming Cheng, Zhongyu Ouyang, Weiyi Wu, Jiang Gui

Pdf | Github Code | Proposed a specialized framework for audio–visual modeling in music understanding, addressing underexplored multimodal interactions, distinctive musical characteristics, and temporal alignment, and introduced annotated rhythmic and source features, with the framework achieving state-of-the-art on Music-AVQA 1.0 and 2.0.
🏷 Multimodal QA 🏷 Representation Learning

[EMNLP 2025] ProtoVQA: An Adaptable Prototypical Framework for Explainable Fine-Grained Visual Question Answering

Oral Presentation Award , 30th EMNLP

Xingjian Diao , Weiyi Wu, Peijun Qing, Keyi Kong, Ming Cheng, Soroush Vosoughi, Jiang Gui

Pdf | Proposed ProtoVQA, an adaptable prototypical VQA framework that learns question-aware prototypes and uses spatially-constrained greedy matching to ground answers in semantically coherent image regions, unifying answering and grounding via a shared backbone and introducing the VLAS metric to quantify visual–linguistic alignment.
🏷 Multimodal QA 🏷 Interpretability

[ACL 2025] Learning Sparsity for Effective and Efficient Music Performance Question Answering

Xingjian Diao , Tianzhen Yang, Chunhui Zhang, Weiyi Wu, Ming Cheng, Jiang Gui

Pdf | Proposed Sparsify, a sparse learning framework for Music Audio-Visual Question Answering that integrates Sparse Masking, Adaptive Sparse Merging, and Sparse Subset Selection to reduce multimodal redundancy, highlight task-critical tokens, and accelerate training convergence, achieving efficiency and accuracy gains across Music-AVQA benchmarks.
🏷 Multimodal QA 🏷 Sparsity Learning

[WACV 2025] FT2TF: First-Person Statement Text-To-Talking Face Generation

Xingjian Diao , Ming Cheng, Wayner Barrios, SouYoung Jin

Pdf | Proposed and developed a one-stage end-to-end text-to-talking face generation pipeline driven by first-person statement text, requiring only visual and textual inputs during inference. Experiments on LRS2 and LRS3 demonstrate state-of-the-art performance, showing its ability to generate realistic talking faces effectively from text inputs.
🏷 AIGC 🏷 Multimodal Alignment

SELECTED COLLABORATIVE PUBLICATIONS

[AACL 2025] Judging the Judges: A Systematic Study of Position Bias in LLM-as-a-Judge

Oral Presentation Award , 14th IJCNLP-AACL

Lin Shi, Chiyu Ma, Wenhua Liang, **Xingjian Diao** , Weicheng Ma, Soroush Vosoughi

Pdf | Proposed a systematic position-bias evaluation framework for LLM-as-a-Judge that integrates Repetition Stability, Position Consistency, and Preference Fairness to measure bias robustness, expose primacy–recency behaviors under prompt permutations, and extend analysis from pairwise to list-wise comparisons on MTBench and DevBench.
🏷 LLM-as-a-Judge

[EMNLP 2025] Knowing More, Acting Better: Hierarchical Representation for Embodied Decision-Making

Chunhui Zhang, Zhongyu Ouyang, **Xingjian Diao** , Zheyuan Liu, Soroush Vosoughi

Pdf | Proposed a hierarchical action probing framework that aggregates layer-wise MLLM representations to enhance spatial grounding, contextual integration, and abstract reasoning for embodied decision-making, achieving substantial gains across language-guided rearrangement tasks.
🏷 MLLM 🏷 Vision–Language–Action

[EMNLP 2025] Assessing and Mitigating Medical Knowledge Drift and Conflicts in Large Language Models

Weiyi Wu, Xinwen Xu, Chongyang Gao, **Xingjian Diao** , Siting Li, Lucas A Salas, Jiang Gui

Pdf | Proposed ConflictMedQA, a guideline-grounded benchmark pairing up-to-date and outdated clinical recommendations, along with ECDA/IKCR metrics to evaluate medical knowledge drift and assess RAG, DPO, and RoD for resolving temporal conflicts in LLMs.
🏷 Medical LLM 🏷 Knowledge Drift 🏷 RAG

[EMNLP 2024] AlphaLoRA: Assigning LoRA Experts Based on Layer Training Quality

Peijun Qing, Chongyang Gao, Yefan Zhou, **Xingjian Diao** , Yaoqing Yang, Soroush Vosoughi

Pdf | Proposed AlphaLoRA, a training-free layer-wise expert allocation strategy for LoRA-MoE that leverages Heavy-Tailed Self-Regularization to quantify layer training quality and assign experts accordingly, reducing redundancy while improving performance across NLP and reasoning benchmarks.
🏷 LLM 🏷 Mixture-of-Experts

[EMBC 2024] GluMarker: A Novel Predictive Modeling of Glycemic Control Through Digital Biomarkers

EMBC NextGen Scholar Award , IEEE 46th EMBC

Ziyi Zhou, Ming Cheng, **Xingjian Diao** , Yanjun Cui, Xiangling Li

Pdf | Proposed GluMarker, a digital biomarker framework that integrates broader daily factors (meals, insulin doses, and CGM-derived metrics) to predict next-day glycemic control and identify key digital biomarkers that reveal how everyday

behaviors shape diabetes management.

🏷 Digital Biomarkers

[Preprint 2025] On The Design Choices of Next Level LLMs

Yijun Tian, Xingjian Diao , Ming Cheng, Chunhui Zhang, Jiang Gui, Soroush Vosoughi, Xiangliang Zhang, Nitesh V. Chawla, Shichao Pei

Pdf | Provided a comprehensive analysis of current LLM design choices across model architecture, attention mechanisms, post-training strategies, optimization techniques, and data selection, identifying key trends and proposing future research directions for next-generation large language models.

🏷 LLM 🏷 Post-Training 🏷 Reinforcement Learning

[Preprint 2025] Tailoring Memory Granularity for Multi-Hop Reasoning over Long Contexts

Peijun Qing, **Xingjian Diao** , Chiyu Ma, Saeed Hassanpour, Soroush Vosoughi

Pdf | Proposed Tailoring Memory Granularity, a reward-guided framework that adaptively composes hybrid memory across multiple granularities, including chunks, triples, atomic facts, and summaries, to enhance large language models in long-context multi-hop reasoning through dynamic query-specific memory selection.

🏷 LLM 🏷 Hybrid Memory 🏷 Multi-hop Reasoning

[Preprint 2025] What Makes a Good Curriculum? Disentangling the Effects of Data Ordering on LLM Mathematical Reasoning

Yaning Jia, Chunhui Zhang, Xingjian Diao , Xiangchi Yuan, Zhongyu Ouyang, Chiyu Ma, Soroush Vosoughi

Pdf | Proposed a unified offline curriculum learning framework that systematically disentangles five curriculum dimensions and isolates the causal impact of data ordering to provide principled guidance on curriculum design.

🏷 LLM 🏷 Curriculum Learning

[Preprint 2025] SPAN: Unlocking Pyramid Representations for Gigapixel Histopathological Images

Weiyi Wu, Xingjian Diao , Chongyang Gao, Xinwen Xu, Siting Li, Jiang Gui

Pdf | Introduced Sparse Pyramid Attention Networks (SPAN) for gigapixel whole-slide pathology, combining spatial-adaptive condensation with context-aware refinement to preserve spatial structure and enable efficient multi-scale tumor detection, classification, and segmentation.

🏷 Digital Pathology 🏷 Whole Slide Image Analysis

[Preprint 2025] Music Performance Audio-Visual Question Answering Requires Specialized Multimodal Designs

Wenhao You, Xingjian Diao , Chunhui Zhang, Keyi Kong, Weiyi Wu, Zhongyu Ouyang, Chiyu Ma, Tingxuan Wu, Noah Wei, Zong Ke, Ming Cheng, Soroush Vosoughi, Jiang Gui

Pdf | Presented the first survey of Music Audio-Visual Question Answering, providing a systematic analysis of how specialized multimodal architectures with spatio-temporal modeling enable reliable reasoning over musical performances.

🏷 Multimodal QA 🏷 Video Understanding

SELECTED SOFTWARE DEVELOPMENT PROJECTS

- **Health Aware Bits (HABits) Lab, Northwestern University** *Sep - Dec 2021*
  Research Assistant Evanston, IL
  Intake Detection Tool with Multiple Classifiers | Github Code
  ◦ Helped develop an approach to detect feeding gestures from the wrist-worn sensor with low inference time and less power consumptiong.
  ◦ Wrote and deployed applications in Android Studio to real devices for evaluation, including training, loading, and testing predefined machine learning algorithms.
  ◦ Implemented the DTW (Dynamic time warping), CNN-LSTM, random forest, SVM, KNN, and Naïve-bayes algorithms using Android Studio.

- **Health Aware Bits (HABits) Lab, Northwestern University** *Mar - Aug 2021*
  Research Assistant Evanston, IL
  Interactive Active Learning Annotation Tool | Github Code
  ◦ Developed and designed an interactive annotation software that uses human-in-the-loop machine learning concepts known as "Active Learning" to label time sequences, with the goal of reducing labeling expenses and addressing challenges faced during the annotation process (used PyQt5, cv2, sklearn, xgboost, numpy, pandas, and pyqtgraph).
  ◦ Added functionalities such as time synchronization, plotting, autoloading of raw data and videos, rewinding video frame by frame, automatically locating queried time sequences, time-sequence-labeling, and labeling results export.
  ◦ Applied the clustered entropy active learning method to query the maximally informative samples.

- **University of Texas Southwestern Medical Center** *May - Aug 2020*
  Software Development Engineer Intern Dallas, TX
  iPADshiny (integrated Protein Array Data management,analysis and visualization tools) | Github Code
  ◦ Developed framework in R shiny for a desktop application that enables biologists to conduct protein-array profiling analysis.
  ◦ Added functionalities for each step of auto-antibody profiling analysis including data import, quality control, normalization procedure, batch correction, and result visualizations with multiple options for every step.

○ Implemented Alone, ANOVA, ComBat, and PCA algorithms for Batch Correction, as well as Scaling, RLM, Quantile, loaess, and VSN algorithms for Normalization, in order to conduct data preprocessing.

TEACHINGS

Graduate Teaching Assistant

- COSC89/189 (Video Understanding), Dartmouth College, Spring 2024
- COSC74/274 (Machine Learning), Dartmouth College, Winter 2024
- COSC61 (Database Systems), Dartmouth College, Summer 2023
- COSC10 (Object Oriented Programming), Dartmouth College, Spring 2023
- COSC62/162 (Applied Cryptography), Dartmouth College, Winter 2023
- COSC10 (Object Oriented Programming), Dartmouth College, Fall 2022

SERVICES

Reviewer / Program Committee Member

- NeurIPS (Annual Conference on Neural Information Processing Systems) 2025
- ICLR (International Conference on Learning Representations) 2025, 2026
- CVPR (Conference on Computer Vision and Pattern Recognition) 2025, 2026
- ICCV (International Conference on Computer Vision) 2025
- ACL (Annual Meeting of the Association for Computational Linguistics) 2025
- ACL (ACL Industry Track) 2025
- EMNLP (Empirical Methods in Natural Language Processing) 2025
- ACMMM (ACM International Conference on Multimedia) 2025
- ACMMM (Datasets Track) 2025
- TMLR (Transactions on Machine Learning Research) 2025
- WACV (IEEE/CVF Winter Conference on Applications of Computer Vision) 2025, 2026
- AISTATS (International Conference on Artificial Intelligence and Statistics) 2026
- EACL (European Chapter of the Association for Computational Linguistics) 2026
- ICASSP (International Conference on Acoustics, Speech & Signal Processing) 2025, 2026
- IUI (ACM International Conference on Intelligent User Interfaces) 2026
- ISBI (IEEE International Symposium on Biomedical Imaging) 2025, 2026
- IJCNN (International Joint Conference on Neural Networks) 2024, 2025
- ICME (IEEE International Conference on Multimedia & Expo) 2024, 2025, 2026
- AVSS (IEEE International Conference on Advanced Video and Signal-Based Systems) 2025
- MINT (NeurIPS Workshop on Foundation Model Interventions) 2024

AWARDS

- **Dartmouth Fellowship,** Dartmouth College. *2022-Present*
- **EMNLP Oral Presentation Award,** 30^{th} EMNLP. *2025*
- **IJCNLP-AACL Oral Presentation Award,** 14^{th} IJCNLP-AACL. *2025*
- **Guarini School of Graduate and Advanced Studies Travel Award,** Dartmouth College. *2025*
- **Biomedical Data Science Travel Award,** Dartmouth College. *2025*
- **IEEE EMBC NextGen Scholar Award,** IEEE EMBC. *2024*